Training-Free Partial-Order Control for Discrete-Diffusion POI Check-in Generation

Anonymous Author(s)

Abstract

Location-based social networks (LBSNs) are a core source of web-scale human-mobility data, yet sharing raw check-in traces is privacy-sensitive, motivating realistic synthetic generation. Beyond statistical realism, many downstream web applications need control over the *semantic order* of activities—e.g., leisure before dining—a sample-level, cross-position requirement that single-step conditioning cannot express. Enforcing such constraints is commonly assumed to cost distributional fidelity; we show this trade-off is largely *separable*. We present PCDG, a *training-free* control stack for a frozen conditional joint temporal + category/POI discrete-diffusion generator. PCDG compiles category partial-order constraints into differentiable energies, projects the category-position logits through a KL-preserving augmented-Lagrangian optimization in the late reverse steps, reduces residual category–POI inconsistency via a category-consistent decoder, and restores spatial realism through a distribution-aware same-category geometry refinement that leaves every constraint metric unchanged *by construction*. Across two cities under an out-of-distribution split and three sampling seeds, PCDG cuts the strict ordering-violation rate from 92.89% to 4.04% on Istanbul and from 89.45% to 11.28% on New York, raising pairwise ordering coverage by 6.5× and 4.2×, while *improving* all six fidelity divergences over the unconstrained backbone on Istanbul and four of six on New York. On Istanbul it surpasses a *fully retrained* partial-order conditional model on all five constraint metrics and five of six fidelity divergences, without retraining anything. Stage-wise ablations localize where fidelity is lost and recovered across the three-stage stack.

CCS Concepts

• **Information systems** → **Location based services**; *Data mining*; • **Computing methodologies** → **Machine learning**; Neural networks; • **Human-centered computing** → *Social networks*.

Keywords

synthetic human mobility, discrete diffusion, partial-order constraints, training-free constrained sampling, POI check-in sequences

1 Introduction

Location-based social networks (LBSNs) produce time-stamped visits to points of interest (POIs), linking human mobility to activity categories such as dining, leisure, and shopping. These traces support POI recommendation, urban analytics, and simulation [9, 25, 40], but releasing raw traces exposes sensitive behavior. Synthetic mobility generation offers an alternative data source [7, 10, 23, 42].

Relevance to the Web. We study a control problem for LBSN data generation: how to generate complete check-in sequences under activity-order requirements while retaining the population statistics used in web-data analysis. Scenario-specific synthetic traces could support recommender stress tests and data augmentation without releasing individual logs. The evaluated contribution is controllable generation; downstream utility and formal privacy protection remain unevaluated.

A scenario may require leisure before dining or sightseeing before shopping. Such precedence couples multiple, non-adjacent positions. A local condition embedding does not ensure that every occurrence of one category precedes every occurrence of another, or that both categories appear. Reordering completed sequences can fix inversions but cannot supply missing activities. Road-topology control [45] addresses geographic structure, whereas constrained discrete diffusion (CDD) [3] provides a general framework for sampling under differentiable constraints. POI generation adds a further difficulty: category tokens and POI identities can disagree, so controlling the former need not control the actual activities of the emitted POIs.

We introduce PCDG, a training-free control stack on the frozen Marionette generator [7]. It adapts CDD's KL-anchored projection to sample-level order and existence energies, with a distance-distribution regularizer on POI logits. Category-consistent decoding then restricts POI support using the sampled category. After diffusion, same-category geometry refinement changes POI identities and coordinates while preserving actual categories, times, and event order. This separates spatial refinement from the fields that determine constraint satisfaction.

Across Istanbul and New York, the stack improves every constraint metric over the retrained PO-CFG baseline. Geometry refinement lowers spatial divergences without changing constraint metrics; the full stack improves all six fidelity JSDs over the unconstrained backbone on Istanbul and four on New York. New York CategoryTransition and G-RANK remain worse. These results identify a recoverable spatial cost in this generator, rather than establishing that every fidelity–control trade-off can be removed.

Contributions.

- We formulate category-precedence control for complete POI check-in sequences using order and existence energies (§3).
- We develop a training-free stack combining logit projection, category-consistent decoding, and geometry refinement within each realized category (§4).
- We separate category-order control from spatial refinement and measure the resulting constraint–fidelity trade-offs, including residual category and popularity errors (§5.4).
- We provide a two-city, three-seed evaluation, projection ablations, and a stage-wise comparison under a fixed selection protocol (§5).

2 Related Work

Human mobility and POI modeling. Classical studies established that human movement is highly regular yet only boundedly predictable [13, 33]. A large body of work predicts the *next* location from history using attentional, spatio-temporal, and graph-enhanced models [9, 25, 40], and diffusion models have recently been applied to next-POI and sequential recommendation [6, 21, 28, 35, 37]. These prediction and ranking objectives differ from generating complete time-stamped visits under sample-level activity-order requirements.

Synthetic trajectory and mobility generation. Data-driven generation spans GAN/Seq2Seq models [2, 18, 20, 39], point-process and neural-ODE generators [23, 42], imitation- and privacy-oriented approaches [10, 17, 27, 29, 34, 36], and, increasingly, diffusion models [4, 12, 22, 30, 38, 41, 44] and LLM-based generators [14, 16, 19, 43]. Marionette supports fine-grained generation using global and partial mobility contexts [7]; we use its generator as a frozen backbone. Related benchmarks examine imitative generation quality [8]. Our target is an explicit relation between activity categories, together with the distributional effects of enforcing it on generated POI identities.

Discrete diffusion. Discrete diffusion extends denoising diffusion [15] to categorical data via structured transition matrices [1], concrete score matching [26], score-entropy [24], and masked formulations [31]; guidance mechanisms steer discrete generation toward conditions [32]. We adopt a discrete-diffusion backbone and intervene at sampling time rather than changing training.

Sampling-time constraint projection. CDD integrates differentiable constraints into discrete-diffusion sampling through KL anchoring and augmented-Lagrangian optimization without retraining [3]. We adapt that projection mechanism to order and existence energies on category positions. The KL anchor penalizes token-level drift from the denoiser; matching aggregate travel-distance and visitation statistics to training references is a separate objective in our stack. We evaluate those terms through projection and stage-wise comparisons rather than treating the KL anchor as a distributional-fidelity guarantee.

Decoding support and spatial refinement. Grammar-constrained decoding restricts language-model outputs to permitted formal structures [11]; Control-DAG uses weighted finite-state automata for lexical, vocabulary, and length control in non-autoregressive generation [5]. Our decoder instead couples category tokens to admissible catalog POIs. Its final-category fallback prevents a universal consistency guarantee. Geographic control is complementary: ControlTraj uses road topology and trip attributes to condition trajectory generation [45]. Our terminal refinement reselects POIs within their realized categories, changing spatial statistics while preserving the fields used by the constraint metrics. The empirical question is how much fidelity this restricted refinement recovers, and which errors remain.

3 Problem Formulation

POI check-in sequences. We model a user's daily behavior on an LBSN as a variable-length sequence of check-in events. A sequence of length n consists of arrival times $t_{1:n}$ (continuous, strictly increasing), activity categories $k_{1:n}$ with $k_i \in \mathcal{C} = \{1, \dots, C\}$, and points of interest $l_{1:n}$ with $l_i \in \mathcal{L}$. Each POI l carries fixed geographic coordinates $\mathrm{gps}(l)$, and each training POI is associated with a category. A per-sequence context $\mathbf{c}$ encodes coarse temporal and behavioral conditions (e.g., time-of-day occupancy indicators). The empirical distribution over such sequences is highly structured: it reflects the regularity and bounded predictability of human mobility [13, 33] while retaining substantial individual variability.

Category partial-order constraints. Many web applications do not merely require a statistically realistic sequence; they require that the *activity semantics* follow an application-specified order, for instance, that leisure precedes dining or that sightseeing precedes shopping. We formalize such requirements as a sample-level *partial order* over categories. Given a reference or user-specified ordering, we encode it as a binary matrix $\mathbf{M} \in \{0, 1\}^{C \times C}$, where $\mathbf{M}[a, b] = 1$ denotes the strict precedence constraint "category a precedes category b": in a satisfying sequence, *every* occurrence of category a must appear before *every* occurrence of category b. When $\mathbf{M}$ is derived from a reference sequence, we set $\mathbf{M}[a, b] = 1$ iff the last position of a precedes the first position of b. A constraint (a, b) is *violated* either by an inversion (some b before some a) or, degenerately, by dropping a or b entirely so that the ordering is vacuously "satisfied." A faithful method must avoid both.

Task. Given a frozen conditional generator of complete check-in sequences and a sample-level partial-order matrix $\mathbf{M}$, we seek to sample

$$(t_{1:n}, k_{1:n}, l_{1:n}) \sim p\big(t_{1:n}, k_{1:n}, l_{1:n} \mid \mathbf{c}, \mathbf{M}\big), \tag{1}$$

such that (i) the generated sequences remain close to the real data distribution in time, space, and activity semantics; and (ii) the category subsequence satisfies the precedence constraints encoded by $\mathbf{M}$. Unlike next-location prediction [9, 25, 40], the task outputs a *complete* time-stamped, categorized, geolocated sequence; unlike single-step conditioning, the constraint couples *multiple, non-adjacent* positions and therefore cannot be enforced by reweighting any single token independently.

Evaluation axes. We evaluate along two axes that are in tension. *Distributional fidelity* is measured by the Jensen–Shannon divergence (JSD) between generated and reference marginals of travel distance, centroid-based Radius, category-transition structure, daily location count, hourly category usage, and a global ranking statistic (G-RANK). *Constraint satisfaction* is measured by ordering-violation rates (OVR), pairwise and category coverage of the required relations, and an overall unsatisfied-relation rate (Unsat); full definitions are given in Appendix D. Reordering can repair inversions but cannot supply missing categories. The goal is to improve constraint satisfaction while retaining distributional fidelity.

4 Method

PCDG controls a frozen generator at sampling time through logit projection, category-consistent POI decoding, and terminal same-category geometry refinement. Figure 1 combines the framework

and a controlled reverse step; ρ, λ, β enter the projection and refinement objectives, rather than forming a separate execution stage.

4.1 Backbone: conditional joint discrete diffusion

We build on a conditional generator that models arrival times with a diffusion point process and the discrete content (categories and POIs) with a conditional discrete diffusion model, following the fine-grained trajectory generator of Deng et al. [7]. The temporal branch samples event times; the discrete branch represents a sequence as a token string $\mathbf{x} = [\textsc{start}, k_{1:n}, \textsc{sep}, l_{1:n}, \textsc{end}]$, with a *category mask* and a *POI mask* identifying the two segments. Categories and POIs use separate absorbing-state transition matrices [1]; a Transformer denoiser with an $\mathbf{x}_0$-parameterization predicts, at each reverse step, per-position logits $\mathbf{z} \in \mathbb{R}^{V \times L}$ over the joint vocabulary of size V across the L positions, from which the posterior $p(\mathbf{x}_{s-1} \mid \mathbf{x}_s)$ and the next discrete state are sampled. PCDG keeps the backbone frozen and changes sampling distributions and terminal POI assignments.

4.2 Partial-order constraints as differentiable energies

Given the sample-level matrix $\mathbf{M}$, we parse it into a set of pairwise constraints $\{(a, b) : \mathbf{M}[a, b] > \frac{1}{2}\}$. To make a constraint differentiable in the logits, we relax the hard categorical choice at each position into a distribution. With temperature T and (optionally) Gumbel noise $\mathbf{g}_l$,

$$\boldsymbol{\pi}_l = \mathrm{softmax}\big((\mathbf{z}_l + \mathbf{g}_l)/T\big) \in \Delta^{V-1}, \qquad l = 1, \ldots, L. \tag{2}$$

We restrict attention to category positions $\mathcal{P}_k$ (via the category mask) and to the category slice of the vocabulary, writing $\pi_l^{\mathrm{cat}}(a)$ for the relaxed probability that position $l \in \mathcal{P}_k$ takes category $a \in C$. For a constraint (a, b) let $P_l^A = \pi_l^{\mathrm{cat}}(a)$ and $P_l^B = \pi_l^{\mathrm{cat}}(b)$. We define two energies:

$$g^{\mathrm{ord}}_{(a,b)} = \sum_{l \in \mathcal{P}_k} \Big(\textstyle\sum_{l' < l} P^B_{l'} \Big) P^A_l, \tag{3}$$

$$g^{\mathrm{ex}}_{(a,b)} = \mathrm{ReLU}\Big(1 - \textstyle\sum_l P^A_l\Big) + \mathrm{ReLU}\Big(1 - \textstyle\sum_l P^B_l\Big). \tag{4}$$

The *order energy* (3) is the expected number of inverted pairs, i.e. the probability mass of a b appearing *before* an a; it is zero iff no b precedes any a. The *existence energy* (4) penalizes an expected count below one for either constrained category, discouraging trivial satisfaction of ordering by *omitting* a or b. Both energies are computed only over category positions; POI logits do not enter Eqs. (3)–(4).

4.3 KL-preserving augmented-Lagrangian projection

At a controlled reverse step, we project the backbone logits to reduce constraint violations while staying close to the model's own distribution. Let $\mathbf{y}^{(0)} = \mathbf{z}$ be the detached backbone logits and $\mathbf{y}$ a trainable copy. We solve, per sequence,

$$\begin{aligned} &\min_{\mathbf{y}} \; D_{\mathrm{KL}}\big(\mathrm{softmax}(\mathbf{y}^{(0)}) \,\big\|\, \mathrm{softmax}(\mathbf{y})\big) \\ &\text{s.t. } g^{\mathrm{ord}}_{(a,b)} \le \tau, \quad g^{\mathrm{ex}}_{(a,b)} \le \tau \quad \forall (a, b), \end{aligned} \tag{5}$$

where the KL term penalizes deviation from the base model; it does not guarantee unchanged fidelity metrics. Its direction is base-to-projected: the frozen distribution is the target argument of the implementation's KL loss. Following the projection paradigm of CDD [3], we optimize the augmented-Lagrangian relaxation of (5): with $\delta^\bullet = \mathrm{ReLU}(g^\bullet - \tau)$,

$$\begin{aligned} \mathcal{L}(\mathbf{y}) = w_{\mathrm{kl}}\, D_{\mathrm{KL}} &+ \sum_{(a,b)} w_{\mathrm{ord}} \Big(\lambda^{\mathrm{ord}} \delta^{\mathrm{ord}} + \tfrac{1}{2} \mu^{\mathrm{ord}} (\delta^{\mathrm{ord}})^2 \Big) \\ &+ \sum_{(a,b)} w_{\mathrm{ex}} \Big(\lambda^{\mathrm{ex}} \delta^{\mathrm{ex}} + \tfrac{1}{2} \mu^{\mathrm{ex}} (\delta^{\mathrm{ex}})^2 \Big). \end{aligned} \tag{6}$$

An *inner* loop takes projected-gradient steps on $\mathbf{y}$ (SGD with learning rate η, gradient-norm clipping, updates restricted to positions of constrained sequences). An *outer* loop evaluates the *hard* violations $\hat{\delta}$ (from the arg max one-hot assignment, not the relaxation) and updates the dual variables,

$$\lambda \leftarrow \lambda + \mu\, \hat{\delta}, \qquad \mu \leftarrow \min(\alpha\, \mu,\ \mu_{\max}) \text{ if } \hat{\delta} > \tau_{\mathrm{tol}}, \tag{7}$$

using a fixed budget without early stopping in the reported runs. Hard violations determine which penalties grow. The order and existence terms act on category positions; distance KL also updates POI logits (§4.4). Hard POI-support masking belongs to the subsequent decoder, not these order/existence energies.

4.4 Distance-distribution regularization

Enforcing category order perturbs the category marginals and, indirectly, the induced spatial statistics. To preserve geographic realism we add an independent distance regularizer that matches the *population* distribution of travel length. From the training split only, we compute the cumulative great-circle (Haversine) route length d of every training sequence with at least two POIs and summarize it as a soft histogram H_{train} over $\log(1 + d)$ with fixed Gaussian bins (the reference is content-hashed and cached, never stored in the model). During projection, for each POI position we fix the top-κ candidate POIs under the backbone and draw P routes with a straight-through Gumbel-softmax (exact distance on the forward pass, a bilinear surrogate on the backward pass), forming a differentiable route-length histogram $H_{\mathrm{proj}}(\mathbf{y})$. We add the term

$$w_{\mathrm{dist}}\, D_{\mathrm{KL}}\big(H_{\mathrm{train}} \,\|\, H_{\mathrm{proj}}(\mathbf{y})\big) \tag{8}$$

to the objective (6). Unlike the order/existence energies, this term acts on *POI-position* logits, and it is a batch-level distribution-matching regularizer, not a hard constraint on each sample. It is applied only to sequences with at least two POIs and uses an independent random stream. Its weight is not tuned per city and is not guaranteed to improve every fidelity metric; we quantify its effect via the No Distance KL ablation (§5).

4.5 Placement in the reverse process and joint sampling

The control is applied in the *late* reverse-diffusion steps, where the predicted categories are informative enough for the energies to be meaningful, specifically during the last K steps and every f steps therein. At each such step, PCDG replaces the backbone category-position logits with their projected values (and, when enabled, the distance-adjusted POI-position logits), and the next discrete state is

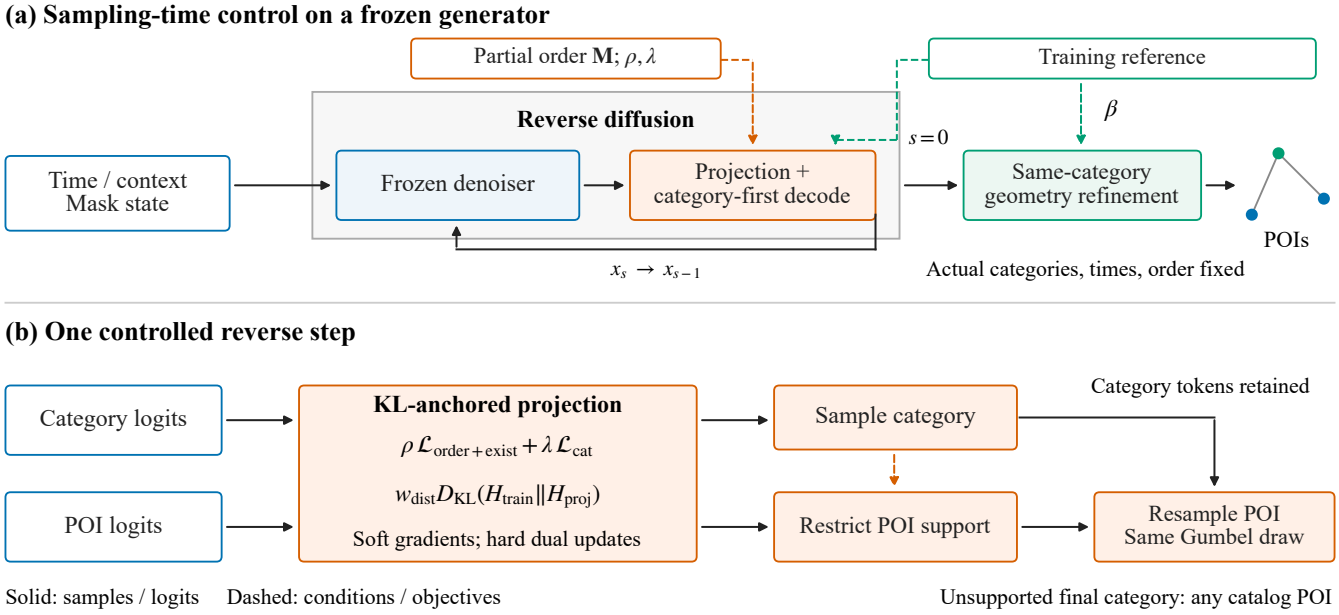


Figure 1: PCDG at sampling time. (a) Projection and decoding operate inside the frozen reverse process; same-category geometry refinement follows it, preserving actual categories, times, and event order. (b) Category energies and distance KL act on different logit positions. The decoder reuses the sampling noise and restricts POI support using the sampled category; unsupported final categories fall back to the catalog. Dashed edges carry conditions or objective inputs.

obtained by the backbone's ordinary joint categorical sampling over the full token sequence, optionally through the category-consistent decoder of §4.6. The temporal branch is unchanged. PCDG adds no model parameters and requires no retraining. Disabling all control stages recovers the unconstrained backbone. Algorithm pseudocode and default hyperparameters are given in Appendix A and Figure 1.

4.6 Category-consistent POI decoding

The category and POI segments are sampled at separate positions. A correctly ordered category-token sequence can therefore yield POIs with different actual categories. Since constraint metrics use $\mathrm{cat}(l)$ from the catalog, projection of category tokens alone does not ensure activity-order control.

The decoder couples the segments within each reverse step. With the backbone's Gumbel draw $\mathbf{g}$, first sample the full token string, fixing the realized category $\hat{c}_l$ aligned with each POI position. Then restrict its POI support and resample using the *same* draw:

$$\hat{x}_l^{\mathrm{cc}} = \arg\max_{v \in S_l} (z_{lv} + g_{lv}), \quad S_l = \{v : \mathrm{cat}(v) = \hat{c}_l\}. \tag{9}$$

For a supported category, logits outside S_l are set to $-\infty$, not a large finite penalty. Category positions are untouched, and the sampler consumes exactly one draw of the original shape per step. At intermediate steps the POI mask token remains admissible; unsupported category positions are unrestricted. At the final step, an unsupported category falls back to *any catalog POI*. Thus the implementation retains a small category–POI mismatch fraction, 0.0034%–0.0219% across the six full-stack runs, rather than eliminating it. The decoder introduces no parameters or logit prior.

4.7 Terminal same-category geometry refinement

The stages so far decide *which* category and *which* POI identity; the coordinates are then whatever the chosen POIs happen to be, and the spatial marginals drift as a result (§5.4). After the reverse process completes, we therefore re-select each POI *within its own realized category*.

For every POI position we build a candidate set of at most κ_g=32 POIs sharing the realized category: the originally sampled POI, the backbone's highest-scoring candidates, the POIs closest to the generated trajectory's own centroid, and, in the distribution-corrected variant, the most frequently visited POIs of the training reference; duplicates are removed, the set is filled by model score, and ties are broken by catalog id so the construction is deterministic. Stage (d) uses the original POI, top-15 by score and nearest-16 to the centroid; stage (e) uses the original, top-11, nearest-12, and the 8 most visited training POIs. Scores come from one frozen, untruncated t=0 evaluation of the backbone; reference GPS from the test split never enters. Writing $\mathbf{w}$ for a relaxed assignment over candidates, we minimize

$$D_{\mathrm{KL}}(q \,\|\, \mathbf{w})\, w_{\mathrm{fid}} + w_{\mathrm{d}}\, \mathrm{JSD}\big(H^{\mathrm{dist}}_{\mathrm{train}}, H^{\mathrm{dist}}(\mathbf{w})\big) + w_{\mathrm{r}}\, \mathrm{JSD}\big(H^{\mathrm{rad}}_{\mathrm{train}}, H^{\mathrm{rad}}(\mathbf{w})\big), \tag{10}$$

where the two soft histograms are differentiable in $\mathbf{w}$ and the proposal q places half of its mass on the original POI and spreads the other half over the candidates by model score, so the fidelity term penalizes drifting away from the backbone's own preference. We optimize Eq. (10) with P_g=8 straight-through relaxed paths, temperature 1, Adam at learning rate 0.05, for a prequalified budget of 200

steps, refreshing an independent geometry noise stream every 20 steps. The output is a *hard* assignment: among the original assignment and the 8 optimized paths we keep the arg-min of the same training-side objective. Empty, single-point and single-candidate positions, and batches in which the base projection never ran, are skipped.

Constraint neutrality. Because the refinement substitutes only POIs of the same realized category, together with their catalog coordinates, and leaves category tokens, the realized category sequence, event times, sequence lengths, event order and all conditions unchanged *per event*, every constraint-relevant statistic is invariant by construction: $\mathrm{OVR}_{\mathrm{strict}}$, $\mathrm{OVR}_{\mathrm{skip}}$, pair and category coverage, Unsat, *Category*, *CategoryTransition* and the category–POI mismatch rate. We verify this numerically on every test run, requiring agreement within 10^{-12} rather than assuming it. The refinement can therefore only move the spatial and popularity statistics while leaving these constraint measurements unchanged.

4.8 Distribution correction

Three scalars govern the balance between constraint strength and distributional fidelity, all referenced to the *training* reference pool only. First, $\rho \in (0, 1]$ scales the constraint term of Eq. (6), so the projection objective becomes

$$w_{\mathrm{kl}}\, D_{\mathrm{KL}} + \rho \sum_{(a,b)} \mathrm{AL}\big(g^{\mathrm{ord}}, g^{\mathrm{ex}}; \lambda^{\bullet}, \mu^{\bullet}\big) + w_{\mathrm{dist}}\, D_{\mathrm{KL}}\big(H_{\mathrm{train}} \| H_{\mathrm{proj}}\big) + \lambda\, \mathcal{L}_{\mathrm{cat}}. \tag{11}$$

Second, $\lambda \geq 0$ weights a category-side term

$$\mathcal{L}_{\mathrm{cat}} = \tfrac{1}{2}\big(\mathrm{JSD}_{\mathrm{hour\times cat}} + \mathrm{JSD}_{\mathrm{transition}}\big), \tag{12}$$

computed from *expected* counts under the relaxed assignment, without Monte Carlo draws or additional randomness. Self-transitions are counted, sequence boundaries excluded, and reference rows without support contribute zero. Third, $\beta \geq 0$ adds to Eq. (10) a POI-side term $\frac{1}{2}(\mathrm{JSD}_{\mathrm{top1000}} + \mathrm{JSD}_{\mathrm{DailyLoc}})$, where the first is an *identity-aligned* top-1000 visitation divergence (the generated and reference top-1000 sets are unioned by POI id before normalization, not compared as sorted popularity curves) and the second a soft expected-*DailyLoc* divergence. Pre-refinement spatial samples are cached independently of β, so the three β settings of a batch are compared at identical upstream randomness.

The full configuration is ρ=0.5, λ=0, β=1, selected on training-side calibration data (Appendix B). The selected value $\lambda = 0$ describes this grid's outcome, not the absence of category-side error.

5 Experiments

We test constraint satisfaction, distributional fidelity, and the contributions of projection, decoding, and geometry refinement.

5.1 Setup

Datasets. New York contains 3,160 training and 2,108 test sequences; Istanbul contains 7,035 and 4,914. New York uses a sequence-level repartition designed to reverse category-precedence preferences between training and test data. This is a distribution shift, not a guarantee of disjoint complete partial-order configurations. Istanbul uses the supplied historical OOD partition; its original split-construction provenance is not fully recoverable. Each test reference supplies a partial-order matrix from actual category first/last positions (§3). Sampling-time methods share frozen city-specific backbones; PO-CFG is separately retrained.

Protocol. Nominal order, existence, token-KL, and distance-KL weights are $(1, 5, 1, 1)$. Projection uses Gumbel temperature 3 and runs every fourth step in the last 40 reverse steps, with 10 outer and 50 inner iterations and no early stopping. Distance KL uses 8 routes, top-32 candidates, and 32 bins. The full stack uses the shared calibration choice $(\rho, \lambda, \beta) = (0.5, 0, 1)$. All reported spreads are sample SDs across seeds 135398, 135399, and 135400 at fixed checkpoints, not training variability or confidence intervals.

Baselines. **JointGen** is the unconstrained backbone. **PostSwap** stably reorders generated events according to **M**. **EnergyGuide** uses a single gradient correction from the constraint energy during sampling, following the guidance paradigm [32]. **PO-CFG** is trained with partial-order conditioning for 1000 epochs and evaluated at classifier-free guidance scale 1; it is the only retrained baseline. Our variants are PCDG-P (projection), PCDG-PD (projection and category-consistent decoding), and PCDG (the full stack).

Metrics. We retain the archived evaluation-v2 protocol for every table. Fidelity uses six JSDs: cumulative travel *Distance*, *Radius*, conditional *CategoryTransition*, distinct daily POI count (*DailyLoc*), hourly *Category* usage, and identity-aligned top-1000 POI visitation (*G-RANK*). Distance is conditional on generated/reference first and last categories matching; the eligible subset can differ by method. Radius is the square root of mean distance to the centroid, not standard radius of gyration. Constraint metrics use the catalog categories of emitted POIs: strict and skip OVR, pair/category coverage, and *Unsat*, the fraction of required relation instances with a missing category or any inversion. Pair coverage requires both categories to occur, not correct ordering. Appendix D specifies denominators and filtering.

5.2 Main results

Table 1 compares all methods on both cities.

Constraint satisfaction. The full PCDG reduces strict OVR from 92.89% to 4.04% on Istanbul and from 89.45% to 11.28% on New York, reductions of 88.85 and 78.17 percentage points. Unsat falls from 95.57% to 7.29% and from 93.64% to 15.86%, while pair coverage increases by 6.5× and 4.2×. All five constraint metrics improve over PO-CFG in both cities. PostSwap reaches zero skip OVR but leaves coverage unchanged, illustrating why skip OVR alone is insufficient. EnergyGuide gives much smaller gains.

Fidelity and residuals. On Istanbul, projection alone degrades Distance (0.1128 → 0.1314) and Radius (0.0226 → 0.0516), while the full stack lowers all six JSDs below JointGen. It also lowers five of six below PO-CFG; its G-RANK remains slightly higher (0.1798 vs. 0.1794). On New York, four of six JSDs improve over JointGen. CategoryTransition (0.0070 → 0.0183) and G-RANK (0.1612 → 0.1914) remain higher, and PO-CFG is best on those two and Category. The observed category-transition gap does not establish that such loss is unavoidable for partial-order control. These comparisons use three-seed means; no significance test is performed.

Table 1: Method comparison (evaluation v2): fidelity JSDs (upper panel) and constraint metrics in percent (lower panel). Cells are mean ± sample SD over three sampling seeds at fixed checkpoints, not confidence intervals. Bold marks the best unrounded city-wise mean, not significance. PCDG-P: projection; PCDG-PD: adds decoding; PCDG: full stack. New York CategoryTransition and G-RANK remain above JointGen.

Istanbul						
Method	Distance ↓	Radius ↓	CategoryTransition ↓	DailyLoc ↓	Category ↓	G-RANK ↓
JointGen	0.1128 ± 0.0087	0.0226 ± 0.0017	0.0403 ± 0.0001	0.0414 ± 0.0009	0.0649 ± 0.0003	0.1922 ± 0.0022
PostSwap	0.0949 ± 0.0127	0.0226 ± 0.0017	0.0361 ± 0.0005	0.0414 ± 0.0009	0.0647 ± 0.0018	0.1922 ± 0.0022
EnergyGuide	0.1156 ± 0.0052	0.0222 ± 0.0025	0.0424 ± 0.0009	0.0392 ± 0.0020	0.0643 ± 0.0016	0.1921 ± 0.0004
PO-CFG (s=1)	0.0890 ± 0.0050	0.0263 ± 0.0020	0.0365 ± 0.0008	0.0444 ± 0.0022	0.0630 ± 0.0007	**0.1794 ± 0.0026**
PCDG-P (projection)	0.1314 ± 0.0035	0.0516 ± 0.0022	**0.0162 ± 0.0006**	0.0368 ± 0.0014	0.0446 ± 0.0020	0.1820 ± 0.0016
PCDG-PD (+decoding)	0.1539 ± 0.0028	0.0892 ± 0.0024	0.0238 ± 0.0011	0.0297 ± 0.0012	0.0423 ± 0.0009	0.1813 ± 0.0018
PCDG (full)	**0.0395 ± 0.0036**	**0.0166 ± 0.0007**	0.0235 ± 0.0011	**0.0249 ± 0.0009**	**0.0423 ± 0.0018**	0.1798 ± 0.0016
New York						
Method	Distance ↓	Radius ↓	CategoryTransition ↓	DailyLoc ↓	Category ↓	G-RANK ↓
JointGen	0.2904 ± 0.0156	0.1104 ± 0.0071	0.0070 ± 0.0008	0.0409 ± 0.0035	0.0184 ± 0.0002	0.1612 ± 0.0013
PostSwap	0.2276 ± 0.0448	0.1104 ± 0.0071	0.0070 ± 0.0007	0.0409 ± 0.0035	0.0168 ± 0.0006	0.1612 ± 0.0013
EnergyGuide	0.2943 ± 0.0060	0.1112 ± 0.0061	0.0071 ± 0.0011	0.0411 ± 0.0029	0.0183 ± 0.0003	0.1622 ± 0.0009
PO-CFG (s=1)	0.1228 ± 0.0160	0.0823 ± 0.0029	**0.0053 ± 0.0018**	0.0499 ± 0.0008	**0.0121 ± 0.0007**	**0.1525 ± 0.0025**
PCDG-P (projection)	0.2210 ± 0.0224	0.1265 ± 0.0051	0.0132 ± 0.0004	0.0351 ± 0.0035	0.0173 ± 0.0013	0.1942 ± 0.0041
PCDG-PD (+decoding)	0.1989 ± 0.0198	0.1681 ± 0.0042	0.0182 ± 0.0012	0.0335 ± 0.0022	0.0178 ± 0.0012	0.1914 ± 0.0059
PCDG (full)	**0.0698 ± 0.0107**	**0.0373 ± 0.0042**	0.0183 ± 0.0005	**0.0301 ± 0.0010**	0.0175 ± 0.0010	0.1914 ± 0.0060

Istanbul					
Method	OVR_strict ↓	OVR_skip ↓	Pair cov. ↑	Category cov. ↑	Unsat ↓
JointGen	92.89 ± 0.08	48.05 ± 0.74	14.63 ± 0.19	43.19 ± 0.20	95.57 ± 0.15
PostSwap	86.10 ± 0.15	**0.00 ± 0.00**	14.63 ± 0.19	43.19 ± 0.20	85.37 ± 0.19
EnergyGuide	90.66 ± 0.06	44.76 ± 0.58	19.17 ± 0.43	47.06 ± 0.25	92.89 ± 0.16
PO-CFG (s=1)	63.45 ± 0.23	19.09 ± 0.54	43.72 ± 0.35	66.00 ± 0.23	69.07 ± 0.45
PCDG-P (projection)	32.80 ± 0.50	9.12 ± 0.10	70.99 ± 0.08	85.36 ± 0.15	36.86 ± 0.37
PCDG-PD (+decoding)	**3.72 ± 0.13**	1.76 ± 0.10	**94.74 ± 0.16**	**97.81 ± 0.07**	**6.90 ± 0.18**
PCDG (full)	4.04 ± 0.20	2.02 ± 0.15	94.68 ± 0.15	97.75 ± 0.10	7.29 ± 0.20
New York					
Method	OVR_strict ↓	OVR_skip ↓	Pair cov. ↑	Category cov. ↑	Unsat ↓
JointGen	89.45 ± 0.20	52.92 ± 0.86	21.13 ± 0.15	46.75 ± 0.12	93.64 ± 0.26
PostSwap	78.22 ± 0.48	**0.00 ± 0.00**	21.13 ± 0.15	46.75 ± 0.12	78.87 ± 0.15
EnergyGuide	89.17 ± 0.14	51.59 ± 0.88	21.17 ± 0.19	47.14 ± 0.14	93.20 ± 0.39
PO-CFG (s=1)	41.91 ± 1.37	8.97 ± 0.88	59.92 ± 1.58	74.99 ± 1.08	49.48 ± 1.20
PCDG-P (projection)	38.51 ± 0.95	14.11 ± 1.22	67.60 ± 1.86	81.54 ± 0.70	42.42 ± 1.84
PCDG-PD (+decoding)	**9.52 ± 1.04**	3.47 ± 0.64	**88.67 ± 1.74**	**93.92 ± 0.68**	**14.09 ± 2.11**
PCDG (full)	11.28 ± 1.05	4.50 ± 0.70	87.87 ± 1.59	93.46 ± 0.63	15.86 ± 1.90

5.3 Ablation of the projection stage

Appendix C reports eight projection-stage variants: No Projection, No Existence, No Order, No KL (both KL terms off), No Distance KL, Fixed Multipliers, No Gumbel, and Full (= PCDG-P).

Removing existence reduces pair coverage to 7.34% on Istanbul and 10.95% on New York, below JointGen's 14.63% and 21.13%. Skip OVR also decreases, to 7.67% and 18.04%: omitting constrained categories evades the order penalty. Removing order retains high coverage (72.53%/69.92%) but raises strict OVR to 63.50%/65.43%. Fixed multipliers increase strict OVR by 9.90/19.94 percentage points, supporting adaptive dual updates at this budget. Distance KL lowers Distance and Radius in both cities: from 0.1441/0.0582 to 0.1314/0.0516 on Istanbul and from 0.2484/0.1430 to 0.2210/0.1265 on New York. Token KL and Gumbel relaxation have small, nonuniform marginal effects; removing Gumbel noise lowers New York Distance to 0.2055.

5.4 Decomposing the fidelity–constraint trade-off

Figure 2 shows the Unsat–Distance plane; the complete stage-wise means are in Table 5, Appendix C.

Projection and decoding. Decoding (b)→(c) reduces strict OVR from 32.80% to 3.72% on Istanbul and from 38.51% to 9.52% on New York. Radius increases from 0.0516 to 0.0892 and from 0.1265 to 0.1681, respectively. The fidelity changes are metric-dependent: category-transition JSD also increases in both cities.

Geometry and distribution correction. At (c)→(d), geometry preserves constraint metrics within 10^{-12}. Distance/Radius fall from 0.1539/0.0892 to 0.0372/0.0159 on Istanbul and from 0.1989/0.1681 to 0.0610/0.0382 on New York. Because actual categories are unchanged, Distance's endpoint filter is also unchanged in this comparison. Spatial fidelity can therefore improve while these constraint

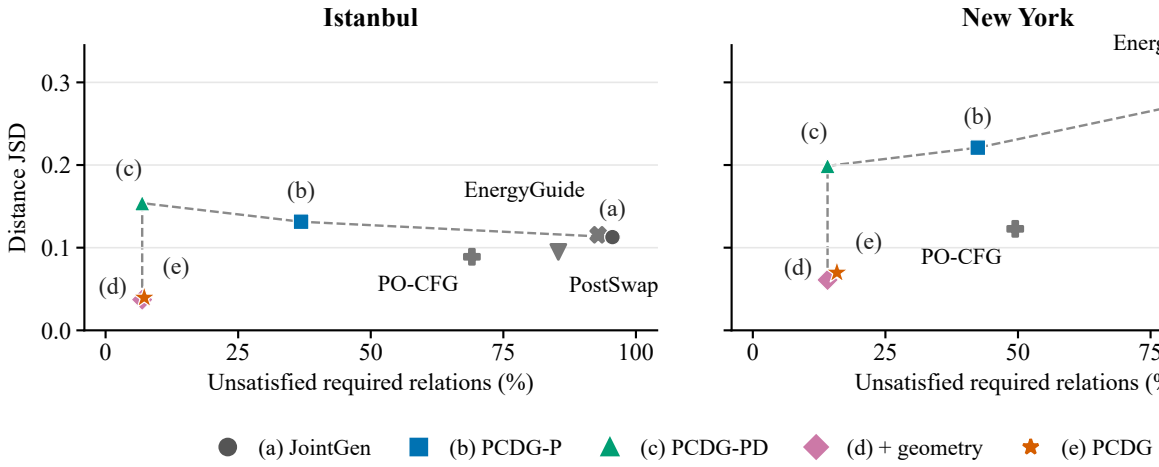


Figure 2: Unsat–Distance trade-off using the evaluation-v2 three-seed means in Tables 1 and 5; lower is better on both axes. Dashed paths connect (a)–(e); shared method/stage points are plotted once. Geometry (c)→(d) lowers Distance JSD without changing Unsat. Transition (d)→(e) changes three factors jointly, not a single regularizer.

statistics stay fixed. New York G-RANK nevertheless rises from 0.1914 to 0.2069. Stage (e) lowers it to 0.1914 and lowers Istanbul G-RANK from 0.1835 to 0.1798, with a modest constraint regression. This transition jointly changes ρ, β, and the candidate pool, so it does not isolate their individual effects. Selection of $\lambda = 0$ identifies a preferred configuration within the evaluated grid, not evidence that category-side fidelity loss is absent.

Selection protocol. The 18 configurations were screened on training-side calibration records already seen by the backbone, not an independent validation set. Ranking minimizes the maximum, then mean, of twelve city–metric JSD/JointGen ratios. The two shortlisted configurations underwent three-seed confirmation; neither passed all distribution gates. The final candidate was chosen from them while retaining the constraint gates, with an explicit override of the failed distribution gates, and evaluated on the test set without reopening selection. Final ranking used v3 re-evaluation; both screening order and final candidate order agree with v2 (Appendix B). No test metric entered hyperparameter selection, candidate construction, or the refinement objective.

5.5 Cost and reproducibility

PCDG requires no retraining but adds projection, decoding, and refinement work. In the earlier geometry qualification, the selected 200-step refinement increased median spatial-sampling time by 3.35–3.98% relative to projection Full, not JointGen. These measurements do not establish overhead for the later full stack. Comparable full-stack/JointGen timing and PO-CFG training GPU-hours are not measured in the archived cost evidence. Appendix E gives measured batch times and their scope.

6 Limitations

The study covers two cities and three sampling seeds at fixed checkpoints, not variation across training runs or a significance test. Istanbul's historical backbone used a different training batch size from New York, and its full training provenance cannot be reconstructed. PO-CFG is evaluated at scale 1, not an extrapolation-tuned optimum. The fixed distance-KL weight is shared across cities; its batch-level objective gives no per-sequence spatial guarantee.

PCDG reduces constraint violations without certifying feasibility. The decoder uses hard POI-support masking but permits a catalog fallback for unsupported final categories (§4.6). New York CategoryTransition (0.0183 vs. 0.0070) and G-RANK (0.1914 vs. 0.1612) remain above JointGen. Hyperparameters were selected on training-side calibration records already seen by the backbone; neither three-seed-confirmed candidate passed all distribution thresholds. Test metrics never entered that selection.

The (d)→(e) comparison changes ρ, β, and candidate-pool composition jointly, precluding single-factor attribution. Stage (d) transfers an already-fixed geometry configuration rather than establishing a separately tuned generalization result. Complex partial-order DAGs, conflict detection, adaptive scheduling, training-time order losses, and downstream recommendation utility remain open. Synthetic output has no formal privacy guarantee here; distributional similarity alone does not establish safe data release.

7 Conclusion

PCDG combines order/existence projection, category-consistent decoding, and same-category geometry refinement for a frozen POI generator. On two cities, it improves all five constraint metrics over retrained PO-CFG. It lowers all six fidelity JSDs versus JointGen on Istanbul and four on New York. Refinement recovers spatial fidelity while preserving constraint-relevant fields; New York's residual category-transition and popularity errors delimit this separation.

Algorithm 1 PCDG: one controlled reverse step

Require: frozen logits $\mathbf{z}$, constraints $\mathbf{M}$, position masks, training references

```
y ← detach(z); initialize duals
for o = 1, ..., 10 do
    for j = 1, ..., 50 do
        Relax y with Gumbel-softmax
        Compute order/existence and distribution terms
        L ← weighted objective in Eq. (11)
        SGD on active rows of y; clip gradients
    end for
    Update duals from hard violations, Eq. (7)
end for
Sample category tokens; decode POIs with shared noise
return next discrete state (Eq. (9) and fallback)
```

A Algorithm and Default Hyper-parameters

Projection uses nominal weights $(w_{\text{ord}}, w_{\text{ex}}, w_{\text{kl}}, w_{\text{dist}}) = (1, 5, 1, 1)$, Gumbel temperature 3, 10 outer and 50 inner iterations, SGD rate $\eta = 1$, and gradient-norm clip 10. Dual variables initialize at 1, penalties grow by 2 up to 10^3, and the violation threshold is $\tau = 0$ with multiplier-update tolerance 10^{-6}. Early stopping is disabled. In the batch implementation, token KL sums over vocabulary and positions for active sequences and divides by the full batch size; order/existence penalties sum over active relations and sequences without that division. The ten controlled steps are $36, 32, \ldots, 4, 0$, totaling 5,000 inner updates per active batch. Distance KL uses 8 routes, top-32 candidates, 32 bins, and path temperature 1.

Geometry uses distance/radius/fidelity weights $(1, 2, 0.1)$, 8 paths, up to 32 candidates, temperature 1, Adam rate 0.05, 200 steps, and noise refresh every 20 steps. The full-stack weights are $(\rho, \lambda, \beta) = (0.5, 0, 1)$. Algorithm 1 covers one controlled reverse step. Geometry acts once after the reverse process; its timing qualification is in Appendix E.

B Distribution-Correction Grid and Selection

We screened $\rho \in \{0.25, 0.5, 1\}$, $\lambda \in \{0, 1\}$, and $\beta \in \{0, 0.1, 1\}$ on 512 training-side conditions per city with one sampling seed. Table 2 retains the archived *v2* screening values. The shared-city ranking minimizes the maximum, then mean, of twelve JSD/JointGen ratios, followed by the parameter tuple to break ties. Upstream samples are cached by city, calibration pool, seed, ρ, and λ, so the three β refinements share upstream randomness and the grid uses a fixed candidate-pool construction.

The top two configurations, $(0.5, 0, 1)$ and $(0.5, 0, 0.1)$, underwent three-seed confirmation. Both retained the constraint gates and met the Distance, Radius, and DailyLoc ceilings, but failed CategoryTransition, Category, and G-RANK ceilings on calibration data. Thus the original all-gates confirmation selected no candidate. Formal sampling subsequently used an explicit override: rank the two confirmed candidates while retaining constraint gates but allowing failed distribution gates.

This final ranking used the archived *v3* deterministic-G-RANK re-evaluation. The 18 screening ranks, two-candidate shortlist, failed confirmation gates, and final candidate order are unchanged under v2. For the selected $(0.5, 0, 1)$, the maximum/mean confirmation ratios are 1.4628/0.8884 in v3 and 1.4628/0.8909 in v2; the runner-up means are 0.9103 and 0.9135, respectively. The selected configuration was frozen before formal test sampling. Test metrics were used only for reporting: eleven of eleven gates hold on Istanbul and nine on New York. The calibration pool was seen during backbone training and is not an independent validation set.

Table 2: Training-side screening (v2): maximum and mean of twelve city–metric JSD/JointGen ratios. * denotes the two candidates sent to three-seed confirmation, not passage of its distribution gates.

Rank	ρ	λ	β	Max	Mean
1*	0.5	0	1	1.7121	0.9148
2*	0.5	0	0.1	1.7121	0.9540
3	0.5	0	0	1.7121	0.9725
4	0.25	0	0	1.7250	0.9477
5	0.25	0	0.1	1.7250	0.9534
6	0.25	0	1	1.7250	0.9604
7	0.5	1	1	1.7515	0.9599
8	0.5	1	0.1	1.7515	0.9739
9	0.5	1	0	1.7515	0.9817
10	0.25	1	0	1.7619	0.9594
11	0.25	1	0.1	1.7619	0.9695
12	0.25	1	1	1.7619	0.9782
13	1	1	1	1.7805	0.9578
14	1	1	0.1	1.7805	0.9951
15	1	1	0	1.7805	1.0046
16	1	0	1	1.7841	0.9501
17	1	0	0.1	1.7841	0.9962
18	1	0	0	1.7841	0.9989

C Projection and Stage-wise Tables

Tables 3 and 4 give the projection ablations; Table 5 reports the full stage-wise comparison. All values follow evaluation v2.

D Metric Definitions

All constraint metrics use actual categories $\text{cat}(l)$ of generated POIs, not the sampled category tokens. For reference sequence i, let R_i contain each distinct-category pair (a, b) satisfying $\max \text{pos}_i(a) < \min \text{pos}_i(b)$. Let $P_i \subseteq R_i$ contain pairs whose two categories both occur in the generated sequence; C_i is the set of categories involved in R_i and $\widehat{C}_i$ the generated category set. For a present pair, $v_i(a, b)$ is the number of inverted occurrence pairs divided by $n_a n_b$.

Ordering and coverage. Strict OVR averages v_i over R_i, assigning 1 to missing-category pairs, then averages over sequences with nonempty R_i. Skip OVR averages v_i over P_i, then over sequences with nonempty P_i. In contrast, coverage and Unsat pool counts:

$$\text{PairCov} = \frac{\sum_i |P_i|}{\sum_i |R_i|}, \tag{13}$$

$$\text{CatCov} = \frac{\sum_i |C_i \cap \widehat{C}_i|}{\sum_i |C_i|}, \tag{14}$$

$$\text{Unsat} = \frac{\sum_i \sum_{(a,b)\in R_i} \mathbf{1}[(a, b) \notin P_i \ \vee \ v_i(a, b) > 0]}{\sum_i |R_i|}. \tag{15}$$

Thus Unsat counts unsatisfied *relations*, not sequences; pair coverage does not require correct order. Empty denominators are undefined, not zero; undefined OVR sequences are omitted from their

Table 3: Ablation of the projection stage: distributional similarity (JSD ↓, evaluation v2). Full here is PCDG-P, i.e. the projection stage alone.

Istanbul						
Variant	Distance ↓	Radius ↓	CategoryTransition ↓	DailyLoc ↓	Category ↓	G-RANK ↓
No Projection	0.1128 ± 0.0087	0.0226 ± 0.0017	0.0403 ± 0.0001	0.0414 ± 0.0009	0.0649 ± 0.0003	0.1922 ± 0.0022
No Existence	**0.1084 ± 0.0114**	**0.0215 ± 0.0027**	0.0389 ± 0.0005	0.0412 ± 0.0021	0.0661 ± 0.0009	0.1924 ± 0.0028
No Order	0.1176 ± 0.0023	0.0486 ± 0.0035	0.0227 ± 0.0008	**0.0358 ± 0.0006**	0.0439 ± 0.0018	**0.1770 ± 0.0040**
No KL	0.1406 ± 0.0117	0.0574 ± 0.0028	0.0164 ± 0.0006	0.0361 ± 0.0012	**0.0429 ± 0.0012**	0.1817 ± 0.0024
Fixed Multipliers	0.1296 ± 0.0093	0.0526 ± 0.0045	0.0185 ± 0.0005	0.0361 ± 0.0009	0.0457 ± 0.0025	0.1841 ± 0.0020
No Gumbel Noise	0.1315 ± 0.0091	0.0524 ± 0.0023	**0.0160 ± 0.0010**	0.0370 ± 0.0005	0.0439 ± 0.0023	0.1831 ± 0.0027
No Distance KL	0.1441 ± 0.0073	0.0582 ± 0.0025	0.0162 ± 0.0010	0.0366 ± 0.0008	0.0437 ± 0.0023	0.1828 ± 0.0026
Full (=PCDG-P)	0.1314 ± 0.0035	0.0516 ± 0.0022	0.0162 ± 0.0006	0.0368 ± 0.0014	0.0446 ± 0.0020	0.1820 ± 0.0016
New York						
Variant	Distance ↓	Radius ↓	CategoryTransition ↓	DailyLoc ↓	Category ↓	G-RANK ↓
No Projection	0.2904 ± 0.0156	0.1104 ± 0.0071	**0.0070 ± 0.0008**	0.0409 ± 0.0035	0.0184 ± 0.0002	0.1612 ± 0.0013
No Existence	0.2439 ± 0.0277	**0.1004 ± 0.0053**	0.0075 ± 0.0008	0.0393 ± 0.0044	0.0189 ± 0.0004	**0.1587 ± 0.0006**
No Order	0.2175 ± 0.0284	0.1196 ± 0.0108	0.0138 ± 0.0003	0.0352 ± 0.0029	0.0177 ± 0.0003	0.1860 ± 0.0048
No KL	0.2415 ± 0.0239	0.1421 ± 0.0046	0.0125 ± 0.0006	0.0352 ± 0.0041	0.0173 ± 0.0013	0.1939 ± 0.0064
Fixed Multipliers	0.2402 ± 0.0487	0.1218 ± 0.0018	0.0127 ± 0.0005	0.0353 ± 0.0041	0.0175 ± 0.0006	0.1891 ± 0.0023
No Gumbel Noise	**0.2055 ± 0.0205**	0.1248 ± 0.0056	0.0130 ± 0.0010	0.0360 ± 0.0036	0.0173 ± 0.0003	0.1918 ± 0.0041
No Distance KL	0.2484 ± 0.0267	0.1430 ± 0.0055	0.0134 ± 0.0011	0.0355 ± 0.0031	**0.0172 ± 0.0013**	0.1937 ± 0.0082
Full (=PCDG-P)	0.2210 ± 0.0224	0.1265 ± 0.0051	0.0132 ± 0.0004	**0.0351 ± 0.0035**	0.0173 ± 0.0013	0.1942 ± 0.0041

Mean ± sample standard deviation over three sampling seeds. Bold denotes the best unrounded mean within each city, not statistical significance. No KL disables both the token-KL anchor and the distance-KL term; No Distance KL disables only the latter. Removing the KL or Gumbel terms does not uniformly degrade fidelity, which we report rather than claiming a benefit for every component.

Table 4: Ablation of the projection stage: constraint satisfaction (%, evaluation v2). The existence energy drives coverage, the order energy drives strict ordering, and the adaptive multipliers beat a fixed penalty.

Istanbul					
Variant	OVR_strict ↓	OVR_skip ↓	Pair cov. ↑	Category cov. ↑	Unsat ↓
No Projection	92.89 ± 0.08	48.05 ± 0.74	14.63 ± 0.19	43.19 ± 0.20	95.57 ± 0.15
No Existence	93.67 ± 0.15	**7.67 ± 0.56**	7.34 ± 0.12	36.92 ± 0.19	93.58 ± 0.17
No Order	63.50 ± 0.76	50.34 ± 0.73	**72.53 ± 0.70**	**86.08 ± 0.17**	70.70 ± 0.73
No KL	32.85 ± 0.52	9.20 ± 0.29	71.11 ± 0.08	85.40 ± 0.16	36.75 ± 0.57
Fixed Multipliers	42.70 ± 0.51	20.97 ± 0.90	71.25 ± 0.89	85.00 ± 0.42	45.64 ± 0.58
No Gumbel Noise	**32.75 ± 0.39**	9.13 ± 0.22	70.97 ± 0.34	85.38 ± 0.19	**36.75 ± 0.25**
No Distance KL	32.96 ± 0.46	9.10 ± 0.16	70.92 ± 0.23	85.29 ± 0.14	36.83 ± 0.24
Full (=PCDG-P)	32.80 ± 0.50	9.12 ± 0.10	70.99 ± 0.08	85.36 ± 0.15	36.86 ± 0.37
New York					
Variant	OVR_strict ↓	OVR_skip ↓	Pair cov. ↑	Category cov. ↑	Unsat ↓
No Projection	89.45 ± 0.20	52.92 ± 0.86	21.13 ± 0.15	46.75 ± 0.12	93.64 ± 0.26
No Existence	90.38 ± 0.47	18.04 ± 1.26	10.95 ± 0.45	39.47 ± 0.12	91.80 ± 0.40
No Order	65.43 ± 0.12	51.24 ± 0.30	**69.92 ± 0.74**	**82.84 ± 0.40**	72.86 ± 0.94
No KL	**38.24 ± 0.96**	13.98 ± 0.82	67.67 ± 1.56	81.57 ± 0.72	**42.40 ± 1.48**
Fixed Multipliers	58.45 ± 0.66	35.14 ± 0.60	63.79 ± 0.76	78.37 ± 0.35	60.67 ± 0.49
No Gumbel Noise	39.11 ± 1.53	13.93 ± 1.34	66.41 ± 1.32	80.95 ± 0.65	43.91 ± 1.59
No Distance KL	38.48 ± 1.04	**13.76 ± 0.72**	67.33 ± 1.86	81.38 ± 0.81	42.77 ± 1.79
Full (=PCDG-P)	38.51 ± 0.95	14.11 ± 1.22	67.60 ± 1.86	81.54 ± 0.70	42.42 ± 1.84

Mean ± sample standard deviation over three sampling seeds. Bold denotes the best unrounded mean within each city. Removing the existence energy collapses coverage *below* the unconstrained backbone, confirming that ordering can otherwise be satisfied by omitting the constrained categories.

respective means. Tables express these rates in percent. The separate diagnostic `token_strict_ovr` uses category tokens. Mismatch rate pools the event-wise disagreements between tokens and catalog categories.

JSD and filtering. For count vectors P, Q, the implementation forms $p = P/\sum P+\epsilon, q = Q/\sum Q+\epsilon, \epsilon = 10^{-14}$, without a second normalization, and evaluates $\frac{1}{2}\sum_j p_j \log(p_j/m_j) + \frac{1}{2}\sum_j q_j \log(q_j/m_j)$, $m = (p+q)/2$, using natural logarithms. A zero-total vector is undefined. Distance, Radius, DailyLoc, and G-RANK retain sequences with at least two GPS points, independently on each side; the two category metrics include shorter sequences. Distance additionally requires the aligned generated/reference first and last categories to match. This subset can vary by method, so the reported Distance is not an unconditional whole-population divergence.

Table 5: Stage-wise comparison, evaluation-v2 three-seed means. SDs for (a), (b), (c), and (e) are in Table 1. Geometry (d) preserves (c)'s constraints within 10^{-12} and uses a previously fixed configuration. Transition (d)→(e) jointly changes ρ $(1 \to 0.5)$, β $(0 \to 1)$, and the same-category candidate pool; it does not isolate a single factor. The calibration grid (Appendix B) holds the pool construction fixed.

	Distributional similarity (JSD ↓)						Constraints (%)		
Stage	Distance	Radius	CatTrans	DailyLoc	Category	G-RANK	OVR_str. ↓	Pair cov. ↑	Unsat ↓
Istanbul									
(a) JointGen (no control)	0.1128	0.0226	0.0403	0.0414	0.0649	0.1922	92.89	14.63	95.57
(b) + order/existence projection	0.1314	0.0516	0.0162	0.0368	0.0446	0.1820	32.80	70.99	36.86
(c) + category-consistent decoding	0.1539	0.0892	0.0238	0.0297	0.0423	0.1813	3.72	94.74	6.90
(d) + geometry refinement	0.0372	0.0159	0.0238	0.0284	0.0423	0.1835	3.72	94.74	6.90
(e) + distribution correction (full)	0.0395	0.0166	0.0235	0.0249	0.0423	0.1798	4.04	94.68	7.29
New York									
(a) JointGen (no control)	0.2904	0.1104	0.0070	0.0409	0.0184	0.1612	89.45	21.13	93.64
(b) + order/existence projection	0.2210	0.1265	0.0132	0.0351	0.0173	0.1942	38.51	67.60	42.42
(c) + category-consistent decoding	0.1989	0.1681	0.0182	0.0335	0.0178	0.1914	9.52	88.67	14.09
(d) + geometry refinement	0.0610	0.0382	0.0182	0.0335	0.0178	0.2069	9.52	88.67	14.09
(e) + distribution correction (full)	0.0698	0.0373	0.0183	0.0301	0.0175	0.1914	11.28	87.87	15.86

Six fidelity statistics. Distance is cumulative consecutive-point Haversine distance (Earth radius 6371 km). With arithmetic latitude/longitude centroid $\bar{x}$, Radius is

$$\text{Radius} = \sqrt{n^{-1} \sum_j d_{\text{Hav}}(x_j, \bar{x})}.$$

Its unit is $\sqrt{\text{km}}$ and it is *not* root-mean-square radius of gyration. DailyLoc counts distinct POI identities per sequence. For these three histograms, edges are exactly $\text{arange}(0, M, M/100)$, where M is the larger observed maximum. This usually produces 99 bins, omits M, drops values beyond the last edge, and renormalizes retained mass. If $M = 0$, the divergence is 0; a zero histogram otherwise yields an undefined value. These historical conventions are preserved for comparability, not proposed as new estimators.

Category averages category-frequency JSDs equally over the 24 hours. CategoryTransition is the equal-weight mean of conditional next-category JSDs over source categories with any outgoing transition on either side; self-transitions are included and sequence boundaries excluded. A row/hour empty on both sides contributes 0; one empty on only one side contributes $\ln 2$. No events or no transitions, respectively, gives an undefined metric. G-RANK retains each side's top-1000 visited POIs (all if fewer), aligns their union by POI identity with zeros for missing retained IDs, and computes JSD after normalizing retained mass. It compares identities, not rank positions.

Deterministic audit. Evaluation v2 uses an unstable frequency sort at the top-1000 cutoff. The v3 audit breaks ties by POI id; identity alignment is present in both versions. Across the 18 formal audit metric sets (six full-stack outputs and twelve reused baseline sets), the largest absolute G-RANK change is 0.0032; other statistics agree within 10^{-12}. The full-stack/backbone G-RANK comparison remains $0.1766 < 0.1892$ on Istanbul and $0.1912 > 0.1628$ on New York under v3. This audit does not replace any v2 table cell.

E Cost

Table 6 reports the earlier selected-geometry qualification, not the later full-stack timing. On each city, ordinary and long cases are eligible 64-sequence batches chosen by lower-median and maximum batchwise generated length, respectively. Each case uses two warmups and five measured repetitions, seed 135398, FP32 with TF32 disabled, and the same CUDA device and input. CUDA synchronization brackets the timer. Ratios divide the two medians, rather than averaging per-trial ratios.

Table 6: Selected 200-step geometry qualification. Median seconds per 64-sequence spatial batch; ratio = Full+Geo / projection Full.

City	Batch	Full	Full+Geo	Ratio
NY	ordinary	27.028	28.063	1.0383
NY	long	33.102	34.211	1.0335
IST	ordinary	24.865	25.856	1.0398
IST	long	27.399	28.371	1.0355

The timer covers spatial sampling and decoding, including fresh scoring, candidate construction, optimization, and discrete selection when geometry is enabled; it excludes temporal generation, model loading, and transfer. The denominator is projection Full (PCDG-P), not JointGen. The selected 200-step configuration uses radius/fidelity weights 2/0.1; an earlier step-budget probe used 1/0.01. All four selected-configuration ratios are below the 1.10 ceiling, with overhead 3.35–3.98%. These measurements cannot be transferred to PCDG-PD plus geometry or the final $(0.5, 0, 1)$ stack. Comparable full-stack/JointGen sampling overhead and PO-CFG training GPU-hours remain unmeasured in the archived cost evidence.

References

[1] Jacob Austin, Daniel Johnson, Jonathan Ho, Daniel Tarlow, and Rianne van den Berg. 2021. Structured Denoising Diffusion Models in Discrete State-Spaces. In *Advances in Neural Information Processing Systems.* arXiv preprint arXiv:2107.03006; NeurIPS 2021.

[2] Chu Cao and Mo Li. 2021. Generating Mobility Trajectories with Retained Data Utility. In *Proceedings of the 27th ACM SIGKDD Conference on Knowledge Discovery and Data Mining.* doi:10.1145/3447548.3467158

[3] Michael Cardei, Jacob K. Christopher, Bhavya Kailkhura, Tom Hartvigsen, and Ferdinando Fioretto. 2025. Constrained Discrete Diffusion. In *Advances in Neural Information Processing Systems 38 (NeurIPS 2025).* doi:10.52202/085713-0415 39th Conference on Neural Information Processing Systems (NeurIPS 2025).

[4] C. Chen, H. Zhang, P. Wang, and F. Lu. 2024. Simulating Human Mobility with a Trajectory Generation Framework Based on Diffusion Model. *International*

Journal of Geographical Information Science (2024). doi:10.1080/13658816.2024.2312199

[5] Jinghong Chen, Weizhe Lin, Jingbiao Mei, and Bill Byrne. 2024. ControlDAG: Constrained Decoding for Non-Autoregressive Directed Acyclic T5 using Weighted Finite State Automata. In *Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 2: Short Papers)*. 508–518. doi:10.18653/v1/2024.naacl-short.42

[6] Jialei Chen, Yuanbo Xu, and Yiheng Jiang. 2025. Unlocking the Power of Diffusion Models in Sequential Recommendation: A Simple and Effective Approach. In *Proceedings of the 31st ACM SIGKDD Conference on Knowledge Discovery and Data Mining*. doi:10.1145/3711896.3737172

[7] Bangchao Deng, Ling Ding, Lianhua Ji, Chunhua Chen, Xin Jing, Bingqing Qu, and Dingqi Yang. 2025. Marionette: Fine-Grained Conditional Generative Modeling of Spatiotemporal Human Trajectory Data Beyond Imitation. In *Proceedings of the 31st ACM SIGKDD Conference on Knowledge Discovery and Data Mining V.2*. 463–473. doi:10.1145/3711896.3737039

[8] Bangchao Deng, Xin Jing, Tianyue Yang, Bingqing Qu, Dingqi Yang, and Philippe Cudre-Mauroux. 2025. Revisiting Synthetic Human Trajectories: Imitative Generation and Benchmarks Beyond Datasaurus. In *Proceedings of the 31st ACM SIGKDD Conference on Knowledge Discovery and Data Mining*. doi:10.1145/3690624.3709180

[9] Jie Feng, Yong Li, Chao Zhang, Funing Sun, Fanchao Meng, Ang Guo, and Depeng Jin. 2018. DeepMove: Predicting Human Mobility with Attentional Recurrent Networks. In *Proceedings of the 2018 World Wide Web Conference*. doi:10.1145/3178876.3186058

[10] Jie Feng, Zeyu Yang, Fengli Xu, Haisu Yu, Mudan Wang, and Yong Li. 2020. Learning to Simulate Human Mobility. In *Proceedings of the 26th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*. doi:10.1145/3394486.3412862

[11] Saibo Geng, Martin Josifoski, Maxime Peyrard, and Robert West. 2023. Grammar-Constrained Decoding for Structured NLP Tasks without Finetuning. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*. doi:10.18653/v1/2023.emnlp-main.674

[12] Letian Gong, Shengnan Guo, Yan Lin, Yichen Liu, Erwen Zheng, Yiwei Shuang, Youfang Lin, Jilin Hu, and Huaiyu Wan. 2025. STCDM: Spatio-Temporal Contrastive Diffusion Model for Check-In Sequence Generation. *IEEE Transactions on Knowledge and Data Engineering* 37, 4 (2025). doi:10.1109/TKDE.2025.3525718

[13] Marta C. González, César A. Hidalgo, and Albert-László Barabási. 2008. Understanding Individual Human Mobility Patterns. *Nature* (2008). doi:10.1038/nature07850

[14] A. Haydari, D. Chen, Z. Lai, et al. 2026. MobilityGPT: Enhanced Human Mobility Modeling With a GPT Model. *IEEE Transactions on Intelligent Transportation Systems* (2026). doi:10.1109/TITS.2025.3626357

[15] Jonathan Ho, Ajay Jain, and Pieter Abbeel. 2020. Denoising Diffusion Probabilistic Models. In *Advances in Neural Information Processing Systems*. arXiv preprint arXiv:2006.11239; NeurIPS 2020.

[16] S. L. Hsu, E. Tung, J. Krumm, et al. 2024. TrajGPT: Controlled Synthetic Trajectory Generation Using a Multitask Transformer-Based Spatiotemporal Model. In *Proceedings of the 32nd ACM International Conference on Advances in Geographic Information Systems*. doi:10.1145/3678717.3691303

[17] Y. Hu, Y. Du, Z. Zhang, et al. 2024. Real-Time Trajectory Synthesis with Local Differential Privacy. In *Proceedings of the IEEE International Conference on Data Engineering*. doi:10.1109/ICDE60146.2024.00137

[18] W. Jiang, W. X. Zhao, J. Wang, and J. Jiang. 2023. Continuous Trajectory Generation Based on Two-Stage GAN. In *Proceedings of the AAAI Conference on Artificial Intelligence*, Vol. 37. doi:10.1609/aaai.v37i4.25557

[19] S. Li, T. Tran, H. Lin, et al. 2025. Geo-Llama: Leveraging LLMs for Human Mobility Trajectory Generation with Constraints. In *Proceedings of the 2025 IEEE International Conference on Mobile Data Management*. doi:10.1109/MDM65600.2025.00023

[20] H. Lin, S. Shaham, Y. Y. Chiang, and C. Shahabi. 2023. Generating Realistic and Representative Trajectories with Mobility Behavior Clustering. In *Proceedings of the 31st ACM International Conference on Advances in Geographic Information Systems*. doi:10.1145/3589132.3625657

[21] Jing Long, Guanhua Ye, Tong Chen, Yang Wang, Meng Wang, and Hongzhi Yin. 2024. Diffusion-Based Cloud-Edge-Device Collaborative Learning for Next POI Recommendations. In *Proceedings of the 30th ACM SIGKDD Conference on Knowledge Discovery and Data Mining*. doi:10.1145/3637528.3671743

[22] Q. Long, C. Rong, T. Li, and Y. Li. 2026. Dynamic Population Distribution Aware Human Trajectory Generation with Diffusion Model. *ACM Transactions on Intelligent Systems and Technology* (2026). doi:10.1145/3787459

[23] Qingyue Long, Huandong Wang, Tong Li, Lisi Huang, Kun Wang, Qiong Wu, Guangyu Li, Yanping Liang, Li Yu, and Yong Li. 2023. Practical Synthetic Human Trajectories Generation Based on Variational Point Processes. In *Proceedings of the 29th ACM SIGKDD Conference on Knowledge Discovery and Data Mining*. doi:10.1145/3580305.3599888

[24] Aaron Lou, Chenlin Meng, and Stefano Ermon. 2024. Discrete Diffusion Modeling by Estimating the Ratios of the Data Distribution. In *Proceedings of the 41st International Conference on Machine Learning*. arXiv preprint arXiv:2310.16834; ICML 2024 Oral.

[25] Yingtao Luo, Qiang Liu, and Zhaocheng Liu. 2021. STAN: Spatio-Temporal Attention Network for Next Location Recommendation. In *Proceedings of the Web Conference 2021*. doi:10.1145/3442381.3449998

[26] Chenlin Meng, Kristy Choi, Jiaming Song, and Stefano Ermon. 2022. Concrete Score Matching: Generalized Score Matching for Discrete Data. In *Advances in Neural Information Processing Systems*, Vol. 35. doi:10.52202/068431-2502

[27] F. Netzler and M. Lienkamp. 2024. Privacy Preserving Human Mobility Generation Using Grid-Based Data and Graph Autoencoders. *ISPRS International Journal of Geo-Information* 13, 7 (2024). doi:10.3390/ijgi13070245

[28] Y. Qin, H. Wu, W. Ju, et al. 2023. A Diffusion Model for POI Recommendation. *ACM Transactions on Information Systems* (2023). doi:10.1145/3624475

[29] J. Rao, S. Gao, and S. Zhu. 2023. CATS: Conditional Adversarial Trajectory Synthesis for Privacy-Preserving Trajectory Data Publication Using Deep Learning Approaches. *International Journal of Geographical Information Science* (2023). doi:10.1080/13658816.2023.2262550

[30] X. Rao, S. Shang, R. Jiang, et al. 2025. Seed: Bridging Sequence and Diffusion Models for Road Trajectory Generation. In *Proceedings of the ACM Web Conference 2025*. doi:10.1145/3696410.3714951

[31] Subham Sahoo, Marianne Arriola, Yair Schiff, Aaron Gokaslan, Edgar Marroquin, Justin Chiu, Alexander Rush, and Volodymyr Kuleshov. 2024. Simple and Effective Masked Diffusion Language Models. In *Advances in Neural Information Processing Systems*, Vol. 37. doi:10.52202/079017-4135

[32] Yair Schiff, Subham Sekhar Sahoo, Hao Phung, Guanghan Wang, Sam Boshar, Hugo Dalla-Torre, Bernardo P. de Almeida, Alexander Rush, Thomas Pierrot, and Volodymyr Kuleshov. 2025. Simple Guidance Mechanisms for Discrete Diffusion Models. In *International Conference on Learning Representations*. arXiv preprint arXiv:2412.10193; ICLR 2025.

[33] Chaoming Song, Zehui Qu, Nicholas Blumm, and Albert-László Barabási. 2010. Limits of Predictability in Human Mobility. *Science* (2010). doi:10.1126/science.1177170

[34] X. Sun, Q. Ye, H. Hu, et al. 2023. Synthesizing Realistic Trajectory Data With Differential Privacy. *IEEE Transactions on Intelligent Transportation Systems* (2023). doi:10.1109/TITS.2023.3241290

[35] D. Wang, Y. Huang, S. Gao, et al. 2025. Generative Next POI Recommendation with Semantic ID. In *Proceedings of the 31st ACM SIGKDD Conference on Knowledge Discovery and Data Mining*. doi:10.1145/3711896.3736981

[36] H. Wang, C. Gao, Y. Wu, et al. 2023. PateGail: A Privacy-Preserving Mobility Trajectory Generator with Imitation Learning. In *Proceedings of the AAAI Conference on Artificial Intelligence*, Vol. 37. doi:10.1609/aaai.v37i12.26700

[37] Ziwei Wang, Jun Zeng, Lin Zhong, Ling Liu, Min Gao, and Junhao Wen. 2024. DSDRec: Next POI Recommendation Using Deep Semantic Extraction and Diffusion Model. *Information Sciences* (2024). doi:10.1016/j.ins.2024.121004

[38] T. Wei, Y. Lin, S. Guo, et al. 2024. Diff-RNTraj: A Structure-Aware Diffusion Model for Road Network-Constrained Trajectory Generation. *IEEE Transactions on Knowledge and Data Engineering* (2024). doi:10.1109/TKDE.2024.3460051

[39] G. Xiong, Z. Li, M. Zhao, et al. 2024. TrajSGAN: A Semantic-Guiding Adversarial Network for Urban Trajectory Generation. *IEEE Transactions on Computational Social Systems* (2024). doi:10.1109/TCSS.2023.3235923

[40] Song Yang, Jiamou Liu, and Kaiqi Zhao. 2022. GETNext: Trajectory Flow Map Enhanced Transformer for Next POI Recommendation. In *Proceedings of the 45th International ACM SIGIR Conference on Research and Development in Information Retrieval*. doi:10.1145/3477495.3531983

[41] Yuan Yuan, Jingtao Ding, Chenyang Shao, Depeng Jin, and Yong Li. 2023. Spatiotemporal Diffusion Point Processes. In *Proceedings of the 29th ACM SIGKDD Conference on Knowledge Discovery and Data Mining*. doi:10.1145/3580305.3599511

[42] Yuan Yuan, Jingtao Ding, Huandong Wang, Depeng Jin, and Yong Li. 2022. Activity Trajectory Generation via Modeling Spatiotemporal Dynamics. In *Proceedings of the 28th ACM SIGKDD Conference on Knowledge Discovery and Data Mining*. doi:10.1145/3534678.3542671

[43] K. Zhang, Y. Pang, Y. Zhang, and Y. Sekimoto. 2024. MobGLM: A Large Language Model for Synthetic Human Mobility Generation. In *Proceedings of the 32nd ACM International Conference on Advances in Geographic Information Systems*. doi:10.1145/3678717.3691311

[44] Yuanshao Zhu, Yongchao Ye, Shiyao Zhang, Xiangyu Zhao, and James J. Q. Yu. 2023. DiffTraj: Generating GPS Trajectory with Diffusion Probabilistic Model. In *Advances in Neural Information Processing Systems*, Vol. 36. doi:10.52202/075280-2843

[45] Yuanshao Zhu, James Jianqiao Yu, Xiangyu Zhao, Qidong Liu, Yongchao Ye, Wei Chen, Zijian Zhang, Xuetao Wei, and Yuxuan Liang. 2024. ControlTraj: Controllable Trajectory Generation with Topology-Constrained Diffusion Model. In *Proceedings of the 30th ACM SIGKDD Conference on Knowledge Discovery and Data Mining*. 4676–4687. doi:10.1145/3637528.3671866